

Fixing flattening's space complexity flaws with flow fusion

Amos Robinson

September 2, 2013

Parallel programming is hard

- ▶ Writing parallel programs by hand is hard and error-prone
- ▶ Nested data parallelism makes it easier for the programmer
- ▶ Compiling nested data parallelism efficiently is hard

Space complexity problem with NDP

Flattening can actually ruin a program's space complexity.
Finding the furthest distance between any points.

```
maxdist :: Vector Point -> Distance
maxdist vs
  = maximum
  $ map maximum
  $ map (\v -> map (distance v) vs) vs
```

Evaluating this naïvely, we end up with an n^2 array in memory:

```
maxdist [p, q, r]
==>
maximum $ map maximum $
[ [distance p p, distance p q, distance p r]
, [distance q p, distance q q, distance q r]
, [distance r p, distance r q, distance r r]]
```

Fusion fixes space complexity

- ▶ Fusion removes intermediate arrays
- ▶ The n^2 array is intermediate
- ▶ So it should be fused away

Short-cut fusion is fragile and relies on inlining

Short-cut fusion:

- ▶ Existing fusion systems are fragile
- ▶ Rely on complicated compiler optimisations
- ▶ Not obvious when fusion will apply

Flow fusion:

- ▶ Flow fusion can fix more cases of space complexity
- ▶ Unknown whether *all* cases
- ▶ Unlikely all cases

Plan

- ▶ Which cases does flow fusion fix
- ▶ Compiler warnings for other cases
- ▶ Prove it

End

end.